

eda_01

July 2, 2025

```
[1]: # CONFIG CELL
from notebook_utils import set_root_directory

set_root_directory()
```

```
[2]: # IMPORTS
import pandas as pd
import os
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
from app.measurement_ingestor import MeasurementIngestor
from app.constants import TARGET_VARIABLES
import dask.dataframe as dd
import dask.dataframe as dd
from dask.diagnostics import ProgressBar
import contextily as ctx
```

```
[3]: # PATH DEFINE
INPUT_DIR = "input_files"
SENSOR_METADATA_PATH = "input_files/sensor_metadata.parquet"

sensor_metadata = pd.read_parquet(SENSOR_METADATA_PATH)
```

```
[4]: class FlatMeasurementIngestor(MeasurementIngestor):
    def transform(self):
        from collections import defaultdict
        import pandas as pd

        time_1h = defaultdict(list)
        time_24h = defaultdict(list)

        # Get all files in directory
        files = [
            os.path.join(self.input_dir, f)
            for f in os.listdir(self.input_dir)
```

```

        if f.endswith((".xlsx", ".xls", ".parquet")) and "sensor_metadata"
↳not in f
    ]

    for file in files:
        file_name = os.path.basename(file)

        try:
            # Handle different file types
            if file.endswith(".parquet"):
                df = pd.read_parquet(file)
            elif file.endswith((".xlsx", ".xls")):
                # Explicitly specify engine for Excel files
                df = pd.read_excel(file, engine="openpyxl") # or 'xlrd'

            # Add your processing logic here
            if "1h" in file_name.lower():
                time_1h[file_name].append(df)
            elif "24h" in file_name.lower():
                time_24h[file_name].append(df)

        except Exception as e:
            print(f"Error processing {file_name}: {str(e)}")
            continue

    # Convert dictionaries to DataFrames
    def concat_dict(d):
        if not d:
            return pd.DataFrame()
        return pd.concat([pd.concat(df_list) for df_list in d.values()])

    return concat_dict(time_1h), concat_dict(time_24h)

```

[5]: # LOADING DATA

```

ingestor = FlatMeasurementIngestor(
    input_dir=INPUT_DIR,
    sensor_metadata=sensor_metadata,
    exclude_depoyzcja=True,
    target_variables=TARGET_VARIABLES,
)

time_1h_df, time_24h_df = ingestor.transform()

# Now these will be DataFrames with proper columns
print("Columns in time_1h_df:", time_1h_df.columns.tolist())

```

```
print("Columns in time_24h_df:", time_24h_df.columns.tolist())
```

Columns in time_1h_df: ['ds', 'unique_id', 'NO2', 'NOx', 'O3', 'SO2', 'C6H6', 'CO', 'PM10', 'PM2.5', 'NO']

Columns in time_24h_df: ['ds', 'unique_id', 'PM10', 'NO2', 'SO2', 'C6H6', 'PM2.5', 'prediction', 'prediction-lo-68', 'prediction-hi-68']

```
[6]: time_1h_df
```

```
[6]:
```

	ds	unique_id	NO2	NOx	O3	SO2	C6H6	CO	PM10	\
0	2000-01-01 01:00:00	10	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
1	2000-01-01 02:00:00	10	6.0	7.0	35.0	NaN	NaN	NaN	NaN	
2	2000-01-01 03:00:00	10	3.0	3.0	42.0	NaN	NaN	NaN	NaN	
3	2000-01-01 04:00:00	10	4.0	4.0	39.0	NaN	NaN	NaN	NaN	
4	2000-01-01 05:00:00	10	3.0	3.0	44.0	NaN	NaN	NaN	NaN	
...	
97197403	2023-12-31 20:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197404	2023-12-31 21:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197405	2023-12-31 22:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197406	2023-12-31 23:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197407	2024-01-01 00:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
...	
97197403	2023-12-31 20:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197404	2023-12-31 21:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197405	2023-12-31 22:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197406	2023-12-31 23:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197407	2024-01-01 00:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
...	
97197403	2023-12-31 20:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197404	2023-12-31 21:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197405	2023-12-31 22:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197406	2023-12-31 23:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
97197407	2024-01-01 00:00:00	1098	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

	PM2.5	NO
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
...
97197403	NaN	NaN
97197404	NaN	NaN
97197405	NaN	NaN
97197406	NaN	NaN
97197407	NaN	NaN

[97197408 rows x 11 columns]

```
[7]: # BASIC STATISTICS
```

```
def basic_stats(df, name):
    print(f"\n=== BASIC STATISTICS - {name} ===")
    print(f"Number of rows: {len(df)}")

    # Handle datetime column
    if "ds" in df.columns:
        print(f>Date range: {df['ds'].min()} - {df['ds'].max()}")
    elif "timestamp" in df.columns:
```

```

print(f"Date range: {df['timestamp'].min()} - {df['timestamp'].max()}")

# Handle ID column
if "unique_id" in df.columns:
    print(f"Unique sensors count: {df['unique_id'].nunique()}")

# Numeric statistics
numeric_cols = df.select_dtypes(include=["number"]).columns
if not numeric_cols.empty:
    print("\nDescriptive statistics for numeric columns:")
    print(df[numeric_cols].describe())
else:
    print("\nNo numeric columns found")

# Categorical statistics (only if such columns exist)
categorical_cols = df.select_dtypes(include=["object", "category"]).columns
if not categorical_cols.empty:
    print("\nDescriptive statistics for categorical columns:")
    print(df[categorical_cols].describe())

# Column types summary
print("\nColumn types summary:")
print(df.dtypes.value_counts())

```

```

[8]: basic_stats(time_1h_df, "Dane 1-godzinnie")
      basic_stats(time_24h_df, "Dane 24-godzinnie")

```

=== BASIC STATISTICS - Dane 1-godzinnie ===

Number of rows: 97197408

Date range: 2000-01-01 01:00:00 - 2024-01-01 00:00:00

Unique sensors count: 462

Descriptive statistics for numeric columns:

	unique_id	NO2	NOx	O3	SO2 \
count	9.719741e+07	2.146812e+07	2.085802e+07	1.380542e+07	1.911445e+07
mean	5.458874e+02	1.773697e+01	2.972659e+01	5.019842e+01	7.623290e+00
std	3.235175e+02	1.647957e+01	4.905571e+01	2.944795e+01	1.174727e+01
min	1.000000e+00	-2.404340e+00	-1.600000e+00	-1.057220e+00	-5.000000e+00
25%	2.420000e+02	6.900000e+00	8.200000e+00	2.800000e+01	2.000660e+00
50%	5.355000e+02	1.271820e+01	1.550000e+01	4.830000e+01	4.115750e+00
75%	8.350000e+02	2.300000e+01	3.050000e+01	6.939440e+01	8.500000e+00
max	1.116000e+03	5.340000e+02	1.877000e+03	5.330000e+02	1.294000e+03

	C6H6	CO	PM10	PM2.5	NO
count	5.456655e+06	1.111020e+07	1.839411e+07	4.222026e+06	5.940571e+06
mean	1.544551e+00	4.529910e-01	2.984016e+01	1.837393e+01	6.580557e+00
std	4.938870e+00	3.882657e-01	3.130964e+01	1.892202e+01	1.916376e+01

min	-5.000000e-02	-1.078600e-01	-6.000000e+00	-3.352276e+00	-9.564340e+00
25%	2.984100e-01	2.467100e-01	1.300000e+01	7.668610e+00	6.808050e-01
50%	7.300000e-01	3.471300e-01	2.140000e+01	1.293211e+01	1.400000e+00
75%	1.665240e+00	5.236900e-01	3.560000e+01	2.218647e+01	3.900000e+00
max	7.980000e+02	1.686000e+01	2.047000e+03	6.442000e+02	9.551610e+02

Column types summary:

float64	9
datetime64[ns]	1
int64	1

Name: count, dtype: int64

=== BASIC STATISTICS - Dane 24-godzinne ===

Number of rows: 8956384

Date range: 2000-01-01 00:00:00 - 2023-12-31 00:00:00

Unique sensors count: 897

Descriptive statistics for numeric columns:

	PM10	NO2	SO2	C6H6 \
count	2.263036e+06	401410.000000	412672.000000	37655.000000
mean	2.857056e+01	19.925634	5.193946	2.413461
std	2.315706e+01	14.132106	9.492338	3.260340
min	0.000000e+00	0.000000	0.000000	0.000000
25%	1.510000e+01	10.200000	1.000000	0.760000
50%	2.200000e+01	17.000000	2.000000	1.500000
75%	3.400000e+01	26.000000	5.800000	3.000000
max	1.098000e+03	293.000000	369.000000	84.940000

	PM2.5	prediction	prediction-lo-68	prediction-hi-68
count	175457.000000	70.000000	70.000000	70.000000
mean	19.264772	23.032094	11.294694	34.769494
std	16.528266	7.948186	7.723040	8.795446
min	0.000000	12.518574	5.638255	19.398894
25%	9.600000	18.763286	6.415980	29.312061
50%	14.300000	20.773875	8.179441	32.751679
75%	22.800000	25.402400	12.653046	40.007629
max	292.000000	53.478715	45.209154	61.748277

Descriptive statistics for categorical columns:

	unique_id
count	8956384
unique	897
top	23
freq	13878

Column types summary:

float64	8
datetime64[ns]	1

object 1
Name: count, dtype: int64

```
[9]: def missing_data_analysis_dask(dask_df, name, sample_size=100000):  
    """  
    Optimized missing data analysis for large datasets using Dask  
    Skips 'ds' and 'unique_id' columns in analysis  
  
    Parameters:  
    - dask_df: Dask DataFrame  
    - name: Analysis name for display  
    - sample_size: Number of rows to sample for visualization  
    """  
    print(f"\n=== MISSING DATA ANALYSIS - {name} ===")  
  
    # 1. Exclude timestamp and ID columns from analysis  
    cols_to_analyze = [col for col in dask_df.columns if col not in ["ds", "  
↪ "unique_id"]]  
    analysis_df = dask_df[cols_to_analyze]  
  
    # 2. Calculate missing values statistics  
    missing = analysis_df.isnull().sum().compute()  
    missing_percent = (missing / len(analysis_df)) * 100  
  
    # Create and display results  
    missing_df = pd.DataFrame(  
        {"Missing Values": missing, "% of Total": missing_percent.round(2)}  
    ).sort_values("Missing Values", ascending=False)  
  
    print(missing_df)  
  
    # 3. Visualization using a sample  
    print(f"\nVisualizing missing data pattern (sample of {sample_size} rows)...  
↪ ")  
  
    # Take a sample for visualization (excluding timestamp/id columns)  
    sample_df = analysis_df.sample(frac=sample_size / len(analysis_df)).  
    ↪ compute()  
  
    plt.figure(figsize=(12, 6))  
    sns.heatmap(sample_df.isnull(), cbar=False, yticklabels=False)  
    plt.title(f"Missing Data Distribution - {name} (Sample)")  
    plt.show()  
  
    # 4. Additional metrics  
    print("\nAdditional Metrics:")  
    print(f"Total rows analyzed: {len(analysis_df):,}")
```

```

print(f"Columns with >50% missing: {sum(missing_percent > 50)}")
print(
    f"Most complete column: {missing_percent.idxmin()} ({missing_percent.
↳min():.2f}% missing)"
)
print(
    f"Least complete column: {missing_percent.idxmax()} ({missing_percent.
↳max():.2f}% missing)"
)

# Usage example:
dask_time_1h = dd.from_pandas(time_1h_df, npartitions=10)
dask_time_24h = dd.from_pandas(time_24h_df, npartitions=10)


missing_data_analysis_dask(dask_time_1h, "1h")
missing_data_analysis_dask(dask_time_24h, "24h")

```

=== MISSING DATA ANALYSIS - 1h ===

	Missing Values	% of Total
PM2.5	92975382	95.66
C6H6	91740753	94.39
NO	91256837	93.89
CO	86087211	88.57
O3	83391985	85.80
PM10	78803295	81.08
SO2	78082958	80.33
NOx	76339392	78.54
NO2	75729292	77.91

Visualizing missing data pattern (sample of 100000 rows)...



Additional Metrics:

Total rows analyzed: 97,197,408

Columns with >50% missing: 9


Most complete column: NO2 (77.91% missing)

Least complete column: PM2.5 (95.66% missing)

=== MISSING DATA ANALYSIS - 24h ===

	Missing Values	% of Total
prediction	8956314	100.00
prediction-lo-68	8956314	100.00
prediction-hi-68	8956314	100.00
C6H6	8918729	99.58
PM2.5	8780927	98.04
NO2	8554974	95.52
SO2	8543712	95.39
PM10	6693348	74.73

Visualizing missing data pattern (sample of 100000 rows)...



Additional Metrics:

Total rows analyzed: 8,956,384

Columns with >50% missing: 8

Most complete column: PM10 (74.73% missing)

Least complete column: prediction (100.00% missing)

```
[10]: def time_analysis_dask(dask_df, name, target_variables, sample_size=100000):
    """
    Optimized time series analysis using Dask for large datasets

    Parameters:
    - dask_df: Dask DataFrame
    - name: Analysis name
    - target_variables: List of columns to analyze
    - sample_size: Sample size for visualizations
    """
    print(f"\n=== TIME SERIES ANALYSIS - {name} (Dask) ===")

    # Create temporary columns for analysis
    dask_df = dask_df.copy()
    if "ds" in dask_df.columns:
        dask_df["date"] = dask_df["ds"].dt.date
        dask_df["hour"] = dask_df["ds"].dt.hour
        dask_df["month"] = dask_df["ds"].dt.month

    # Take a sample for visualization
```

```

with ProgressBar():
    print(f"\nSampling {sample_size} rows for visualization...")
    df_sample = dask_df.sample(frac=sample_size / len(dask_df)).compute()

# Visualization
plt.figure(figsize=(18, 12))

# 1. Monthly trends (computed in parallel)
with ProgressBar():
    print("\nCalculating monthly averages...")
    monthly = dask_df.groupby("month")[target_variables].mean().compute()

plt.subplot(2, 2, 1)
monthly.plot(kind="bar", ax=plt.gca())
plt.title(f"Monthly Averages - {name}")
plt.ylabel("Concentration (µg/m³)")

# 2. Hourly trends (computed in parallel)
with ProgressBar():
    print("\nCalculating hourly averages...")
    hourly = dask_df.groupby("hour")[target_variables].mean().compute()

plt.subplot(2, 2, 2)
hourly.plot(kind="line", ax=plt.gca())
plt.title(f"Hourly Averages - {name}")
plt.ylabel("Concentration (µg/m³)")

# 3. Boxplots (on sample)
plt.subplot(2, 2, 3)
sns.boxplot(data=df_sample[target_variables])
plt.title(f"Value Distribution - {name} (Sample)")
plt.xticks(rotation=45)
plt.ylabel("Concentration (µg/m³)")

# 4. Daily trends (computed in parallel)
with ProgressBar():
    print("\nCalculating daily trends...")
    daily = dask_df.groupby("date")[target_variables].mean().compute()

plt.subplot(2, 2, 4)
daily.plot(ax=plt.gca())
plt.title(f"Daily Trends - {name}")
plt.ylabel("Concentration (µg/m³)")

plt.tight_layout()
plt.show()

```

```

return dask_df.drop(["date", "hour", "month"], axis=1)

# Convert to Dask DataFrames
dask_time_1h = dd.from_pandas(time_1h_df, npartitions=10)
dask_time_24h = dd.from_pandas(time_24h_df, npartitions=10)

# Define target variables
TARGET_1H = ["NO2", "NOx", "O3", "SO2", "C6H6", "CO", "PM10", "PM2.5", "NO"]
TARGET_24H = ["NO2", "PM10", "SO2", "C6H6", "PM2.5"]

# Execute analysis
print("\nStarting analysis for 1-hour data...")
dask_time_1h = time_analysis_dask(dask_time_1h, "1-hour data", TARGET_1H)

print("\nStarting analysis for 24-hour data...")
dask_time_24h = time_analysis_dask(dask_time_24h, "24-hour data", TARGET_24H)

```

Starting analysis for 1-hour data...

=== TIME SERIES ANALYSIS - 1-hour data (Dask) ===

Sampling 100000 rows for visualization...

```

[#####] | 100% Completed | 101.72 ms
[#####] | 100% Completed | 12.05 s

```

Calculating monthly averages...

```

[#####] | 100% Completed | 3.53 sms

```

Calculating hourly averages...

```

[#####] | 100% Completed | 1.93 sms

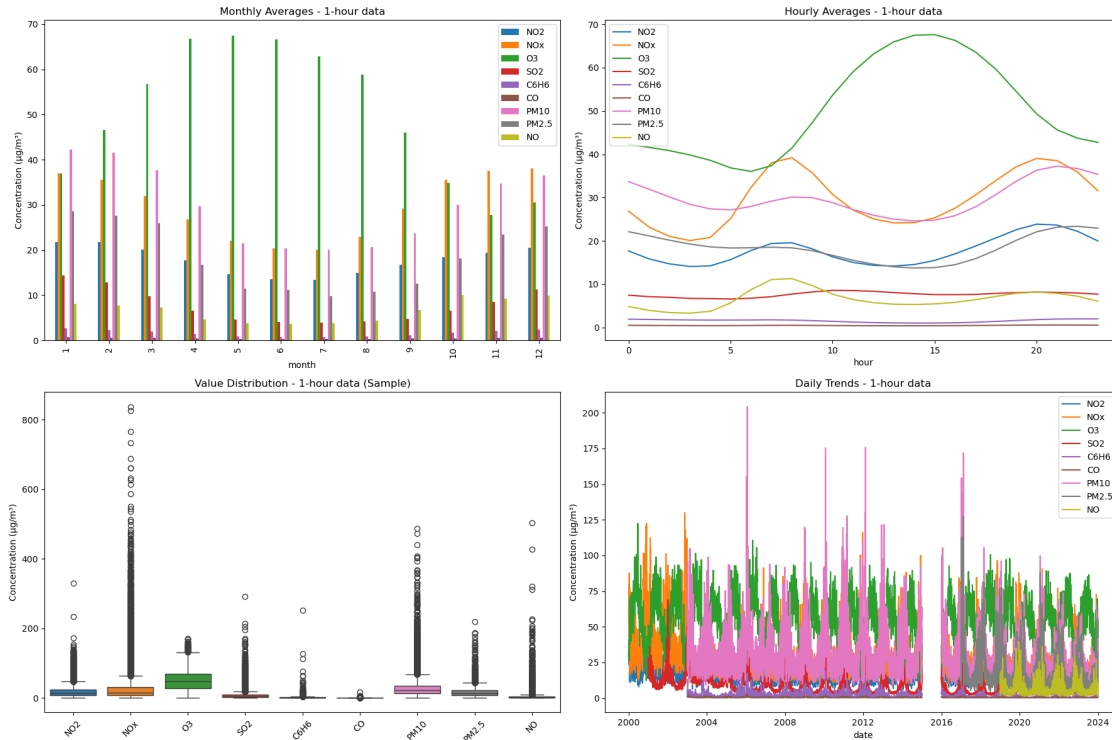
```

Calculating daily trends...

```

[#####] | 100% Completed | 16.96 s

```



Starting analysis for 24-hour data...

=== TIME SERIES ANALYSIS - 24-hour data (Dask) ===

Sampling 100000 rows for visualization...

[#####] | 100% Completed | 105.12 ms

[#####] | 100% Completed | 1.24 sms

Calculating monthly averages...

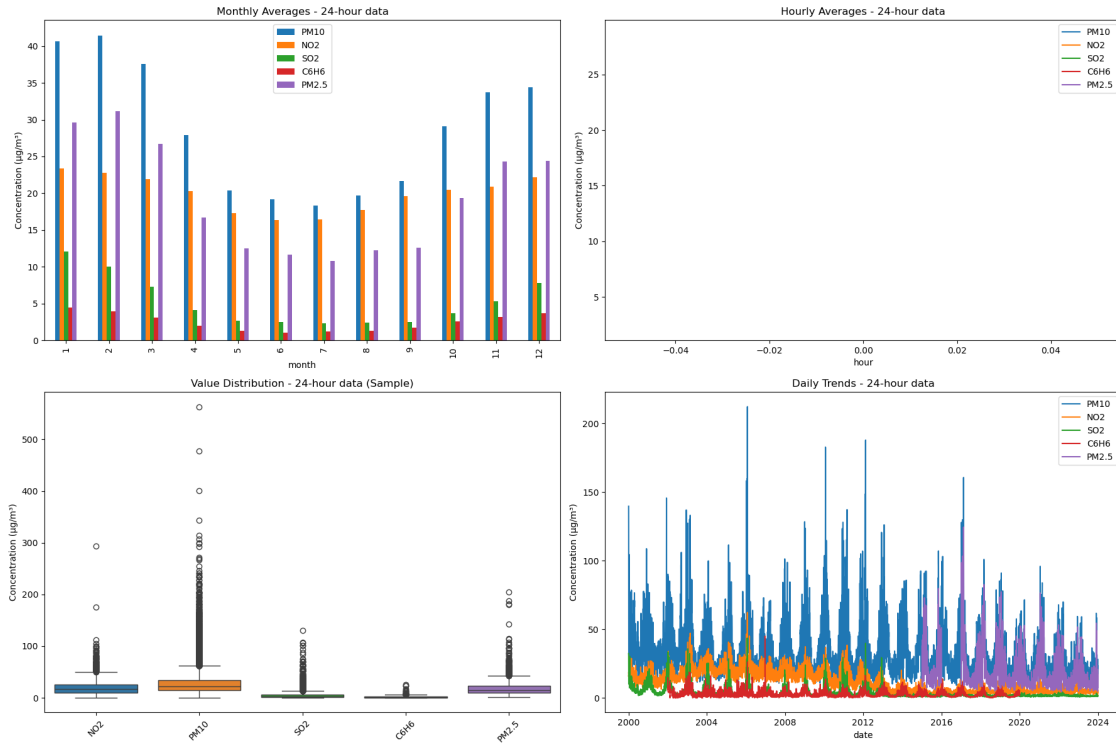
[#####] | 100% Completed | 205.72 ms

Calculating hourly averages...

[#####] | 100% Completed | 204.48 ms

Calculating daily trends...

[#####] | 100% Completed | 1.74 sms



```
[22]: # SPATIAL ANALYSIS
def spatial_analysis_dask(dask_df, sensor_metadata, target_var="PM2.5",
    sample_size=1000):
    """
    Spatial analysis with Dask optimization and Poland map background

    Parameters:
    - dask_df: Dask DataFrame with measurement data
    - sensor_metadata: Pandas DataFrame with sensor locations
    - target_var: Variable to visualize (default: PM2.5)
    - sample_size: Number of points to sample for visualization
    """
    sensor_metadata = sensor_metadata.copy()
    dask_df = dask_df.copy()

    if not all(col in sensor_metadata.columns for col in ["latitude",
    "longitude"]):
        print("Skipping spatial analysis - missing coordinates")
        return

    print("\n=== SPATIAL ANALYSIS WITH POLAND MAP (Dask) ===")

    with ProgressBar():
```

```

# Merge with metadata
sensor_metadata["sensor_id"] = sensor_metadata["sensor_id"].astype(str)
dask_df["unique_id"] = dask_df["unique_id"].astype(str)
merged = dd.merge(
    dask_df[["unique_id", target_var]],
    sensor_metadata[["sensor_id", "latitude", "longitude"]],
    left_on="unique_id",
    right_on="sensor_id",
).compute()

sample = merged.sample(min(sample_size, len(merged)))

# Create figure and axis
plt.figure(figsize=(14, 10))
ax = plt.gca()

# Calculate plot boundaries (Poland approximate coordinates)
min_lon, max_lon = 14.0, 24.5
min_lat, max_lat = 49.0, 55.0

# Filter points within Poland boundaries
sample = sample[
    (sample["longitude"].between(min_lon, max_lon))
    & (sample["latitude"].between(min_lat, max_lat))
]

# Create scatter plot
scatter = ax.scatter(
    x=sample["longitude"],
    y=sample["latitude"],
    c=sample[target_var],
    s=sample[target_var].rank(pct=True) * 100 + 20,
    alpha=0.7,
    cmap="viridis",
    edgecolors="w",
    linewidths=0.5,
)

# Add map background
ctx.add_basemap(
    ax,
    crs="EPSG:4326", # WGS84 coordinate system
    source=ctx.providers.OpenStreetMap.Mapnik,
    zoom=7,
)

# Customize plot

```

```

plt.title(f"{target_var} Spatial Distribution in Poland (Sample)",
↪fontsize=14, pad=20)
cbar = plt.colorbar(scatter, ax=ax, shrink=0.7)
cbar.set_label(f"{target_var} Concentration (µg/m³)", fontsize=12)

ax.set_xlabel("Longitude", fontsize=12)
ax.set_ylabel("Latitude", fontsize=12)
ax.set_xlim(min_lon, max_lon)
ax.set_ylim(min_lat, max_lat)

plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

[23]: # CORRELATION ANALYSIS

```

def correlation_analysis_dask(dask_df, name, target_variables):
    """
    Correlation analysis with Dask optimization

    Parameters:
    - dask_df: Dask DataFrame
    - name: Analysis name
    - target_variables: List of columns for correlation
    """
    print(f"\n=== CORRELATION ANALYSIS - {name} (Dask) ===")

    with ProgressBar():
        # Compute correlation matrix in parallel
        corr_matrix = dask_df[target_variables].corr().compute()

    plt.figure(figsize=(12, 8))
    sns.heatmap(
        corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", center=0, vmin=-1,
↪vmax=1, square=True
    )
    plt.title(f"Correlation Matrix - {name}")
    plt.tight_layout()
    plt.show()

```

[24]:

```

spatial_analysis_dask(dask_time_1h, sensor_metadata, "PM2.5")
spatial_analysis_dask(dask_time_24h, sensor_metadata, "PM2.5")

correlation_analysis_dask(dask_time_1h, "1h", TARGET_1H)
correlation_analysis_dask(dask_time_24h, "24h", TARGET_24H)

```

=== SPATIAL ANALYSIS WITH POLAND MAP (Dask) ===

c:\Users\Mambo\.conda\envs\ispies\Lib\site-packages\dask\dataframe\multi.py:169:

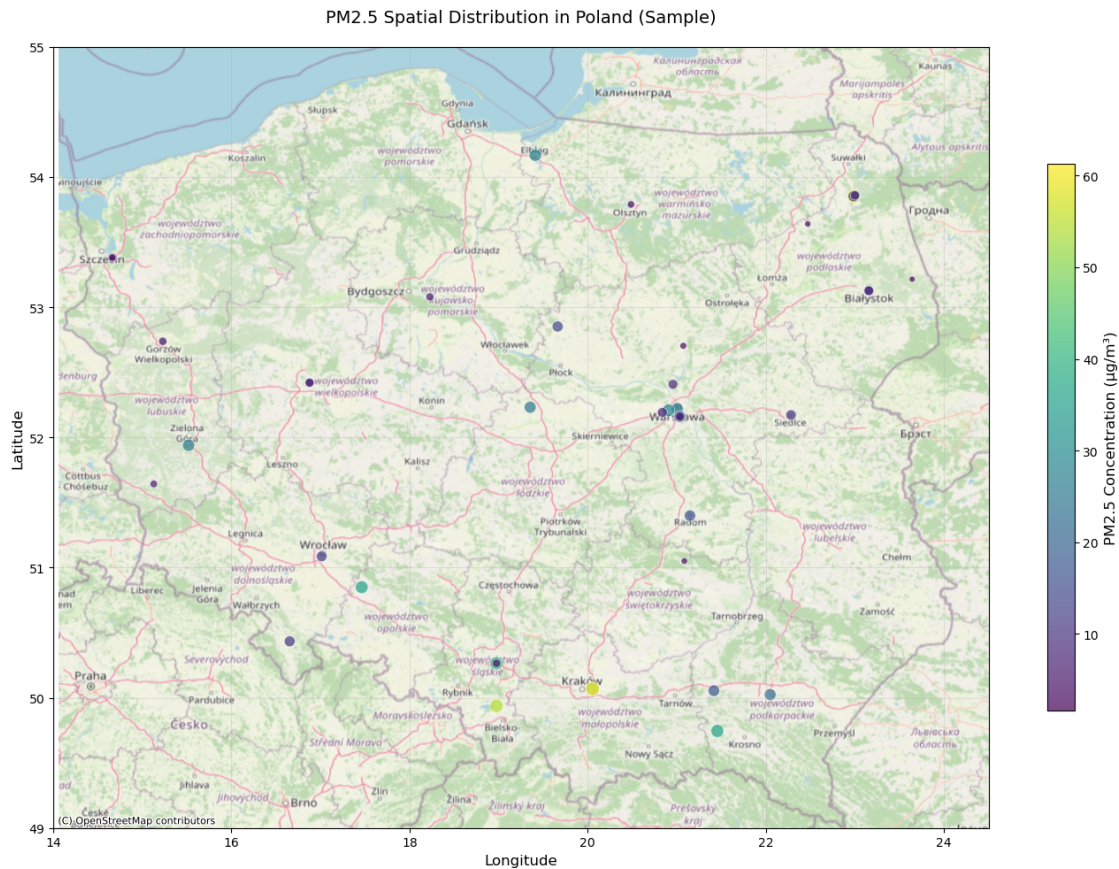
UserWarning: Merging dataframes with merge column data type mismatches:

```
+-----+-----+-----+
| Merge columns          | left dtype | right dtype |
+-----+-----+-----+
| ('unique_id', 'sensor_id') | object     | string      |
+-----+-----+-----+
```

Cast dtypes explicitly to avoid unexpected results.

warnings.warn(

[#####] | 100% Completed | 21.19 s



=== SPATIAL ANALYSIS WITH POLAND MAP (Dask) ===

c:\Users\Mambo\.conda\envs\ispies\Lib\site-packages\dask\dataframe\multi.py:169:

UserWarning: Merging dataframes with merge column data type mismatches:

```
+-----+-----+-----+
| Merge columns          | left dtype | right dtype |
+-----+-----+-----+
```

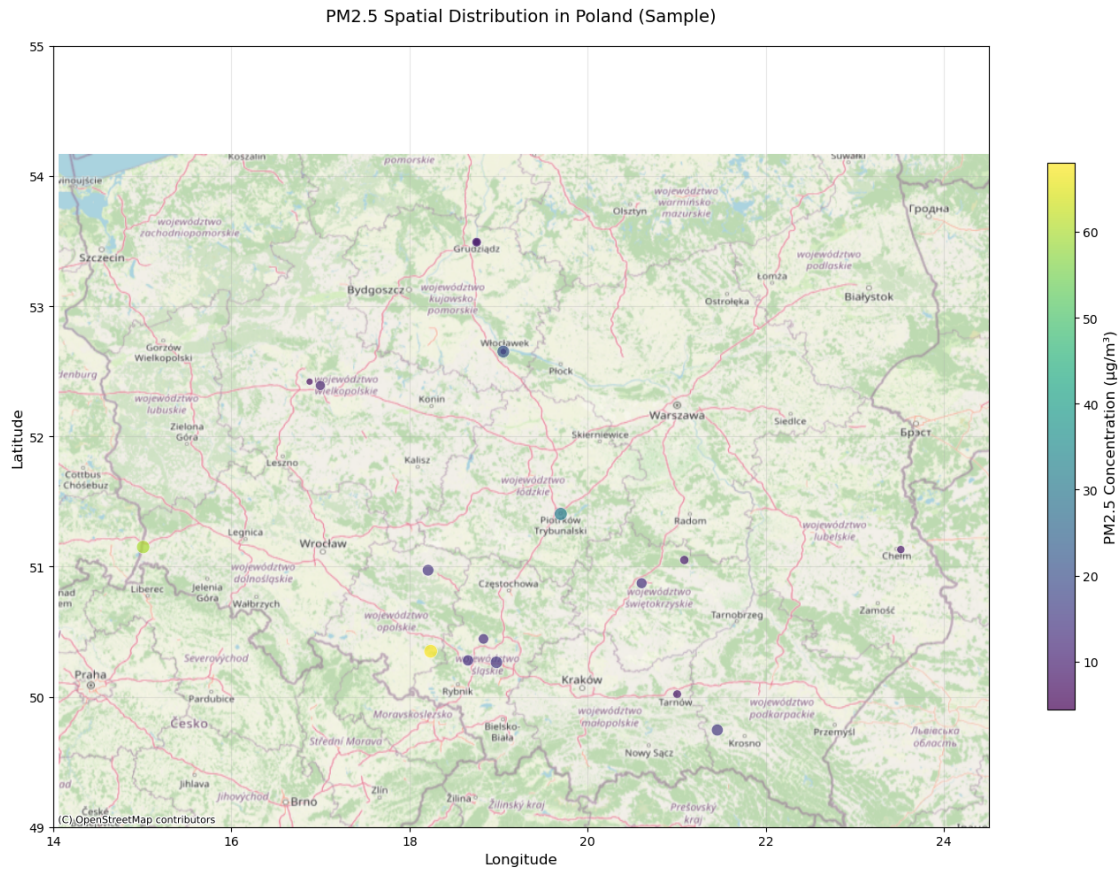


```
| ('unique_id', 'sensor_id') | object      | string      |
+-----+-----+-----+-----+-----+-----+
```

Cast dtypes explicitly to avoid unexpected results.


```
warnings.warn(
```

```
[#####] | 100% Completed | 2.17 sms
```




=== CORRELATION ANALYSIS - 1h (Dask) ===

```
[#####] | 100% Completed | 22.90 s
```



=== CORRELATION ANALYSIS - 24h (Dask) ===

[#####] | 100% Completed | 814.69 ms



```
[25]: # OUTLIER ANALYSIS
def outlier_analysis_dask(dask_df, name, target_variables, n_partitions=10):
    """
    Outlier analysis using Dask with boxplots and statistical summary

    Parameters:
    - dask_df: Dask DataFrame
    - name: Analysis name
    - target_variables: List of columns to analyze
    - n_partitions: Number of partitions for Dask
    """
    print(f"\n=== OUTLIER ANALYSIS - {name} (Dask) ===")

    # Convert to Dask if not already
    if not isinstance(dask_df, dd.DataFrame):
        dask_df = dd.from_pandas(dask_df, npartitions=n_partitions)
```

```

# Calculate basic statistics in parallel
with ProgressBar():
    stats = dask_df[target_variables].describe(percentiles=[0.01, 0.99]).
↪compute()

print("\nExtreme Value Statistics:")
display(stats.loc[["min", "1%", "99%", "max"]])

# Calculate IQR-based outliers in parallel
with ProgressBar():
    q1 = dask_df[target_variables].quantile(0.25).compute()
    q3 = dask_df[target_variables].quantile(0.75).compute()
    iqr = q3 - q1

    lower_bound = q1 - 1.5 * iqr
    upper_bound = q3 + 1.5 * iqr

    outlier_counts = {}
    for col in target_variables:
        outliers = (
            dask_df[(dask_df[col] < lower_bound[col]) | (dask_df[col] >
↪upper_bound[col]))[col]
            .count()
            .compute()
        )
        outlier_counts[col] = outliers

print("\nIQR-based Outlier Counts (1.5*IQR rule):")
outlier_df = pd.DataFrame.from_dict(outlier_counts, orient="index",
↪columns=["Outlier Count"])
outlier_df["% of Total"] = (outlier_df["Outlier Count"] / len(dask_df)) *
↪100
display(outlier_df.round(2))

# Sample data for visualization
sample_size = min(100000, len(dask_df) // 10)
sample_df = dask_df.sample(frac=sample_size / len(dask_df)).compute()

# Visualization
plt.figure(figsize=(15, 6))
sns.boxplot(data=sample_df[target_variables])
plt.title(f"Outlier Distribution - {name} (Sample of {sample_size:,} rows)")
plt.xticks(rotation=45)
plt.ylabel("Concentration (µg/m³)")
plt.show()

return {

```

```

        "stats": stats,
        "outlier_counts": outlier_df,
        "iqr_bounds": pd.DataFrame({"lower": lower_bound, "upper":
↪upper_bound}),
    }

outlier_results_1h = outlier_analysis_dask(dask_time_1h, "1-hour data",
↪TARGET_1H)

outlier_results_24h = outlier_analysis_dask(dask_time_24h, "24-hour data",
↪TARGET_24H)

```

=== OUTLIER ANALYSIS - 1-hour data (Dask) ===

[#####] | 100% Completed | 6.13 sms

Extreme Value Statistics:

	N02	NOx	O3	S02	C6H6	CO	PM10 \
min	-2.404340	-1.6	-1.05722	-5.0	-0.05000	-0.107860	-6.000000
1%	1.586367	2.2	2.10000	0.5	0.07072	0.157688	3.380894
99%	88.500000	325.0	131.00000	65.6	23.80000	2.173000	188.760000
max	534.000000	1877.0	533.00000	1294.0	798.00000	16.860000	2047.000000

	PM2.5	NO
min	-3.352276	-9.564340
1%	2.471120	0.110269
99%	119.086260	134.209580
max	644.200000	955.161000

[#####] | 100% Completed | 655.87 ms

[#####] | 100% Completed | 698.82 ms

[#####] | 100% Completed | 102.69 ms

[#####] | 100% Completed | 102.32 ms

[#####] | 100% Completed | 102.47 ms

[#####] | 100% Completed | 102.32 ms

[#####] | 100% Completed | 102.99 ms

[#####] | 100% Completed | 102.34 ms

[#####] | 100% Completed | 102.91 ms

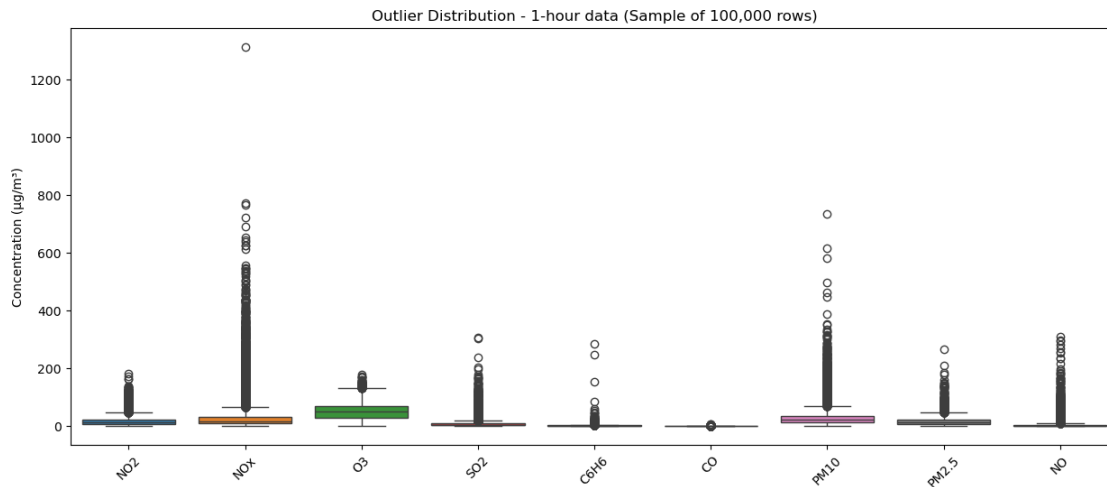
[#####] | 100% Completed | 102.29 ms

[#####] | 100% Completed | 102.24 ms

IQR-based Outlier Counts (1.5*IQR rule):

	Outlier Count	% of Total
N02	1002567	1.03
NOx	1698707	1.75

O3	54763	0.06
S02	1097371	1.13
C6H6	309497	0.32
CO	552334	0.57
PM10	1062389	1.09
PM2.5	232513	0.24
NO	679708	0.70



=== OUTLIER ANALYSIS - 24-hour data (Dask) ===

[#####] | 100% Completed | 413.62 ms

Extreme Value Statistics:

	NO2	PM10	S02	C6H6	PM2.5
min	0.0	0.000	0.0	0.00	0.000
1%	3.0	6.400	0.4	1.00	6.148
99%	87.0	164.189	51.0	20.00	105.450
max	293.0	1098.000	369.0	84.94	292.000

[#####] | 100% Completed | 104.03 ms

[#####] | 100% Completed | 103.40 ms

[#####] | 100% Completed | 102.94 ms

[#####] | 100% Completed | 102.67 ms

[#####] | 100% Completed | 102.49 ms

[#####] | 100% Completed | 102.64 ms

[#####] | 100% Completed | 102.54 ms

IQR-based Outlier Counts (1.5*IQR rule):

	Outlier Count	% of Total
NO2	8461	0.09

PM10	81704	0.91
SO2	29130	0.33
C6H6	2053	0.02
PM2.5	6863	0.08

